

CODE: Sentiment Classification - Impact of RNN

Architecture

```
import numpy as np

import tensorflow_datasets as tfds
import tensorflow as tf

data, info = tfds.load('imdb_reviews', with_info=True, as_supervised=True)
train, test = data['train'], data['test']

[> Downloading and preparing dataset imdb_reviews/plain_text/1.0.0 (download: 80.23
DI Completed...:    0/0 [00:00<?, ? url/s]

DI Size...:    0/0 [00:00<?, ? MiB/s]

Shuffling and writing examples to /root/tensorflow_datasets/imdb_reviews/plain_t
100%                               24999/25000 [00:00<00:00, 105187.70 examples/s]
Shuffling and writing examples to /root/tensorflow_datasets/imdb_reviews/plain_t
100%                               24999/25000 [00:00<00:00, 117584.58 examples/s]
Shuffling and writing examples to /root/tensorflow_datasets/imdb_reviews/plain_t
100%                               49999/50000 [00:00<00:00, 133195.65 examples/s]
WARNING:absl:Dataset is using deprecated text encoder API which will be removed :
Dataset imdb_reviews downloaded and prepared to /root/tensorflow_datasets/imdb_r

VOCAB_SIZE = 10000
BATCH_SIZE = 64
embedding_dim=16
max_length = 120
trunc_type= 'post'
UNK="<UNK>"

## Processing the data

train_input = [] #X_train
test_input = [] #X_test
train_labels = [] #y_train
test_labels = [] #y_test

for i,j in train:
    train_input.append(str(i.numpy()))
```

```
train_labels.append(j.numpy())

for i,j in test:
    test_input.append(str(i.numpy()))
    test_labels.append(j.numpy())

train_labels = np.array(train_labels)
test_labels = np.array(test_labels)

from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences

tokenizer = Tokenizer(num_words = VOCAB_SIZE, oov_token=UNK)
tokenizer.fit_on_texts(train_input)
word_index = tokenizer.word_index

seq = tokenizer.texts_to_sequences(train_input)
padding = pad_sequences(seq, maxlen=max_length, truncating = trunc_type)
test_seq = tokenizer.texts_to_sequences(test_input)
test_padded = pad_sequences(test_seq, maxlen=max_length)

print("train : ", len(seq[0]))
print("test : ", len(test_seq[0]))

    train : 118
    test : 173

## Vanilla RNN Model

model = tf.keras.Sequential([
    tf.keras.layers.Embedding(VOCAB_SIZE, embedding_dim, input_length=max_length),
    tf.keras.layers.SimpleRNN(64),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
model.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 120, 16)	160000
simple_rnn_1 (SimpleRNN)	(None, 64)	5184
dense_2 (Dense)	(None, 10)	650
dense_3 (Dense)	(None, 1)	11

```
=====
Total params: 165,845
Trainable params: 165,845
Non-trainable params: 0
=====
```

```
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
history=model.fit(padding, train_labels, epochs=10, validation_data = (test_padded, te
```

```
Epoch 1/10
782/782 [=====] - 140s 174ms/step - loss: 0.6951 - accu:
Epoch 2/10
782/782 [=====] - 134s 172ms/step - loss: 0.6949 - accu:
Epoch 3/10
782/782 [=====] - 134s 172ms/step - loss: 0.6942 - accu:
Epoch 4/10
782/782 [=====] - 137s 175ms/step - loss: 0.6914 - accu:
Epoch 5/10
782/782 [=====] - 133s 170ms/step - loss: 0.6677 - accu:
Epoch 6/10
782/782 [=====] - 134s 172ms/step - loss: 0.5931 - accu:
Epoch 7/10
782/782 [=====] - 135s 173ms/step - loss: 0.6056 - accu:
Epoch 8/10
782/782 [=====] - 134s 171ms/step - loss: 0.5453 - accu:
Epoch 9/10
782/782 [=====] - 133s 170ms/step - loss: 0.5762 - accu:
Epoch 10/10
782/782 [=====] - 132s 169ms/step - loss: 0.5765 - accu:
```

```
from sklearn import metrics
from sklearn.metrics import precision_score, recall_score
```

```
predicted_1 = model.predict(test_padded)
```

```
pred_1 = np.zeros(len(predicted_1))
for i, score in enumerate(predicted_1):
    if score > 0.5:
        pred_1[i] = 1
    else:
        predicted_1[i] = 0
```

```
print(metrics.classification_report(test_labels, pred_1))
print("Precision = ",precision_score(test_labels, pred_1))
print("Recall = ", recall_score(test_labels, pred_1))
```

	precision	recall	f1-score	support
0	0.58	0.41	0.48	12500
1	0.55	0.70	0.61	12500
accuracy			0.56	25000

macro avg	0.56	0.56	0.55	25000
weighted avg	0.56	0.56	0.55	25000

```
Precision = 0.5451332920024798
Recall = 0.70344
```

```
import matplotlib.pyplot as plt
```

```
def plot_graphs(history, input):
    plt.plot(history.history[input])
    plt.plot(history.history['val_'+input])
    plt.xlabel("Epochs")
    plt.ylabel(input)
    plt.legend([input, 'val_'+input])
    plt.show()
```

```
plot_graphs(history, 'accuracy')
plot_graphs(history, 'loss')
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-2-924d1a230a6c> in <module>()
      9     plt.show()
     10
----> 11 plot_graphs(history, 'accuracy')
     12 plot_graphs(history, 'loss')

NameError: name 'history' is not defined
```

SEARCH STACK OVERFLOW

```
## GRU Model
```

```
model_3 = tf.keras.Sequential([
    tf.keras.layers.Embedding(VOCAB_SIZE, embedding_dim, input_length=max_length),
    tf.keras.layers.Bidirectional(tf.keras.layers.GRU(64)),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
model_3.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
embedding_3 (Embedding)	(None, 120, 16)	160000
bidirectional (Bidirectional)	(None, 128)	31488
dense_4 (Dense)	(None, 32)	4128

dense_5 (Dense) (None, 1) 33

```
=====
Total params: 195,649
Trainable params: 195,649
Non-trainable params: 0
=====
```

```
model_3.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
history = model_3.fit(padding, train_labels, epochs=10, validation_data = (test_padded,
```

```
Epoch 1/10
782/782 [=====] - 59s 69ms/step - loss: 0.4952 - accuracy: 0.1250
Epoch 2/10
782/782 [=====] - 53s 68ms/step - loss: 0.2936 - accuracy: 0.2500
Epoch 3/10
782/782 [=====] - 53s 68ms/step - loss: 0.2275 - accuracy: 0.3750
Epoch 4/10
782/782 [=====] - 52s 67ms/step - loss: 0.1679 - accuracy: 0.5000
Epoch 5/10
782/782 [=====] - 53s 67ms/step - loss: 0.1124 - accuracy: 0.6250
Epoch 6/10
782/782 [=====] - 53s 68ms/step - loss: 0.0711 - accuracy: 0.7500
Epoch 7/10
782/782 [=====] - 53s 68ms/step - loss: 0.0401 - accuracy: 0.8750
Epoch 8/10
782/782 [=====] - 53s 68ms/step - loss: 0.0272 - accuracy: 0.9000
Epoch 9/10
782/782 [=====] - 54s 69ms/step - loss: 0.0236 - accuracy: 0.9250
Epoch 10/10
782/782 [=====] - 54s 69ms/step - loss: 0.0146 - accuracy: 0.9500
```

```
from sklearn import metrics
from sklearn.metrics import precision_score, recall_score
```

```
predicted_3 = model_3.predict(test_padded)
```

```
pred_3 = np.zeros(len(predicted_3))
for i, score in enumerate(predicted_3):
    if score > 0.5:
        pred_3[i] = 1
    else:
        pred_3[i] = 0
```

```
print(metrics.classification_report(test_labels, pred_3))
print("Precision = ", precision_score(test_labels, pred_3))
print("Recall = ", recall_score(test_labels, pred_3))
```

	precision	recall	f1-score	support
0	0.81	0.83	0.82	12500
1	0.83	0.81	0.82	12500



accuracy			0.82	25000
macro avg	0.82	0.82	0.82	25000
weighted avg	0.82	0.82	0.82	25000

Precision = 0.826363414037662

Recall = 0.81096

```
plot_graphs(history, 'accuracy')
```

```
plot_graphs(history, 'loss')
```



```
## LSTM Model
```

```
model_2 = tf.keras.Sequential([
    tf.keras.layers.Embedding(VOCAB_SIZE, embedding_dim, input_length=max_length),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64)),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
model_2.summary()
```

Model: "sequential_3"

Layer (type)	Output Shape	Param #
embedding_4 (Embedding)	(None, 120, 16)	160000
bidirectional_1 (Bidirectional)	(None, 128)	41472
dense_6 (Dense)	(None, 32)	4128
dense_7 (Dense)	(None, 1)	33
Total params: 205,633		
Trainable params: 205,633		
Non-trainable params: 0		

```
model_2.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
history = model_2.fit(padding, train_labels, epochs=10, validation_data = (test_padded,
```

```
Epoch 1/10
782/782 [=====] - 61s 73ms/step - loss: 0.4469 - accuracy: 0.0000
Epoch 2/10
782/782 [=====] - 56s 71ms/step - loss: 0.2918 - accuracy: 0.0000
Epoch 3/10
782/782 [=====] - 55s 71ms/step - loss: 0.2358 - accuracy: 0.0000
Epoch 4/10
782/782 [=====] - 55s 70ms/step - loss: 0.1899 - accuracy: 0.0000
Epoch 5/10
782/782 [=====] - 55s 71ms/step - loss: 0.1535 - accuracy: 0.0000
Epoch 6/10
782/782 [=====] - 55s 71ms/step - loss: 0.1118 - accuracy: 0.0000
Epoch 7/10
782/782 [=====] - 55s 71ms/step - loss: 0.0831 - accuracy: 0.0000
Epoch 8/10
782/782 [=====] - 55s 71ms/step - loss: 0.0628 - accuracy: 0.0000
Epoch 9/10
782/782 [=====] - 56s 71ms/step - loss: 0.0453 - accuracy: 0.0000
Epoch 10/10
782/782 [=====] - 56s 72ms/step - loss: 0.0332 - accuracy: 0.0000
```

```
from sklearn import metrics
from sklearn.metrics import precision_score, recall_score
```

```
predicted_2 = model_2.predict(test_padded)
```

```
pred_2 = np.zeros(len(predicted_2))
for i, score in enumerate(predicted_2):
    if score > 0.5:
        pred_2[i] = 1
    else:
        pred_2[i] = 0
```



```

for i in range(0,split_size):
    small_test.append(test_labels[sorted_list_ind[i]])

for i in range(split_size, split_size*2):
    med_test.append(test_labels[sorted_list_ind[i]])

for i in range(split_size*2, len(test_labels)):
    long_test.append(test_labels[sorted_list_ind[i]])

print(len(small_test))
print(len(med_test))
print(len(long_test))

```

```

8333
8333
8334

```

```

## Model - 1 : Vanilla RNN

```

```

small_pred_1 = []
med_pred_1 = []
long_pred_1 = []

for i in range(0,split_size):
    small_pred_1.append(pred_1[sorted_list_ind[i]])

for i in range(split_size, split_size*2):
    med_pred_1.append(pred_1[sorted_list_ind[i]])

for i in range(split_size*2, len(test_labels)):
    long_pred_1.append(pred_1[sorted_list_ind[i]])

```

```

## Small Input

```

```

print("====RNN: SMALL INPUT====")
print("Precision = ",precision_score(small_test, small_pred_1))
print("Recall = ", recall_score(small_test, small_pred_1))

```

```

## Medium Input

```

```

print("====RNN: MEDIUM INPUT====")
print("Precision = ",precision_score(med_test, med_pred_1))
print("Recall = ", recall_score(med_test, med_pred_1))

```

```

## Long Input

```

```

print("====RNN: LONG INPUT====")
print("Precision = ",precision_score(long_test, long_pred_1))
print("Recall = ", recall_score(long_test, long_pred_1))

```

```

====RNN: SMALL INPUT====
Precision = 0.5822124071471592

```

```

Recall = 0.6708304418228083
=====RNN: MEDIUM INPUT=====
Precision = 0.5303951367781155
Recall = 0.6880236569738788
=====RNN: LONG INPUT=====
Precision = 0.5269328802039083
Recall = 0.7528526341344987

```

```
## Model - 2 : LSTM
```

```

small_pred_2 = []
med_pred_2 = []
long_pred_2 = []

```

```

for i in range(0,split_size):
    small_pred_2.append(pred_2[sorted_list_ind[i]])

```

```

for i in range(split_size, split_size*2):
    med_pred_2.append(pred_2[sorted_list_ind[i]])

```

```

for i in range(split_size*2, len(test_labels)):
    long_pred_2.append(pred_2[sorted_list_ind[i]])

```

```
## Small Input
```

```

print("=====RNN: SMALL INPUT=====")
print("Precision = ",precision_score(small_test, small_pred_2))
print("Recall = ", recall_score(small_test, small_pred_2))

```

```
## Medium Input
```

```

print("=====RNN: MEDIUM INPUT=====")
print("Precision = ",precision_score(med_test, med_pred_2))
print("Recall = ", recall_score(med_test, med_pred_2))

```

```
## Long Input
```

```

print("=====RNN: LONG INPUT=====")
print("Precision = ",precision_score(long_test, long_pred_2))
print("Recall = ", recall_score(long_test, long_pred_2))

```

```

=====RNN: SMALL INPUT=====
Precision = 0.8214360477287
Recall = 0.9077029840388618
=====RNN: MEDIUM INPUT=====
Precision = 0.7567911714770797
Recall = 0.8787580088713652
=====RNN: LONG INPUT=====
Precision = 0.7309113300492611
Recall = 0.864530225782957

```

```
## Model - 3 : GRU
```

```
small_pred_3 = []
```

```

small_pred_3 = []
med_pred_3 = []
long_pred_3 = []

for i in range(0,split_size):
    small_pred_3.append(pred_3[sorted_list_ind[i]])

for i in range(split_size, split_size*2):
    med_pred_3.append(pred_3[sorted_list_ind[i]])

for i in range(split_size*2, len(test_labels)):
    long_pred_3.append(pred_3[sorted_list_ind[i]])


## Small Input
print("====RNN: SMALL INPUT====")
print("Precision = ",precision_score(small_test, small_pred_3))
print("Recall = ", recall_score(small_test, small_pred_3))

## Medium Input
print("====RNN: MEDIUM INPUT====")
print("Precision = ",precision_score(med_test, med_pred_3))
print("Recall = ", recall_score(med_test, med_pred_3))

## Long Input
print("====RNN: LONG INPUT====")
print("Precision = ",precision_score(long_test, long_pred_3))
print("Recall = ", recall_score(long_test, long_pred_3))

====RNN: SMALL INPUT====
Precision =  0.868986454927604
Recall =  0.8607448531112654
====RNN: MEDIUM INPUT====
Precision =  0.8142352347299344
Recall =  0.7949728930507639
====RNN: LONG INPUT====
Precision =  0.7929405915983098
Recall =  0.7744598203447439

```

 0s completed at 8:29 AM

