

CODE: Sentiment Classification - Impact of Pretrained Word Embedding

```

1 import numpy as np
2
3 import tensorflow_datasets as tfds
4 import tensorflow as tf
5
6 data, info = tfds.load('imdb_reviews', with_info=True, as_supervised=True)
7 train, test = data['train'], data['test']

```

Downloading and preparing dataset imdb_reviews/plain_text/1.0.0 (download: 80.23

DI Completed...: 100% 1/1 [00:05<00:00, 5.32s/ url]

DI Size...: 100% 80/80 [00:05<00:00, 23.93 MiB/s]

Shuffling and writing examples to /root/tensorflow_datasets/imdb_reviews/plain_t
100% 24999/25000 [00:00<00:00, 116024.35 examples/s]

Shuffling and writing examples to /root/tensorflow_datasets/imdb_reviews/plain_t
100% 24999/25000 [00:00<00:00, 118347.06 examples/s]

Shuffling and writing examples to /root/tensorflow_datasets/imdb_reviews/plain_t
100% 49999/50000 [00:00<00:00, 133471.40 examples/s]

WARNING:absl:Dataset is using deprecated text encoder API which will be removed
Dataset imdb_reviews downloaded and prepared to /root/tensorflow_datasets/imdb_r

```

1 VOCAB_SIZE = 10000
2 BATCH_SIZE = 64
3 embedding_dim=100
4 max_length = 120
5 trunc_type= 'post'
6 UNK="<UNK>"

```

```

1 ## Processing the data
2
3 train_input = [] #X_train
4 test_input = [] #X_test
5 train_labels = [] #y_train
6 test_labels = [] #y_test
7
8 for i,j in train:
9     train_input.append(str(i.numpy()))

```

```

10     train_labels.append(j.numpy())
11
12 for i,j in test:
13     test_input.append(str(i.numpy()))
14     test_labels.append(j.numpy())
15
16
17 train_labels = np.array(train_labels)
18 test_labels = np.array(test_labels)

```

```
1 print(test_input[1])
```

e they concern his moral integrity and we are never quite sure whether it remain:

```

1 from tensorflow.keras.preprocessing.text import Tokenizer
2 from tensorflow.keras.preprocessing.sequence import pad_sequences
3
4 tokenizer = Tokenizer(num_words = VOCAB_SIZE, oov_token=UNK)
5 tokenizer.fit_on_texts(train_input)
6 word_index = tokenizer.word_index

1 seq = tokenizer.texts_to_sequences(train_input)
2 padding = pad_sequences(seq, maxlen=max_length, truncating = trunc_type)
3 test_seq = tokenizer.texts_to_sequences(test_input)
4 test_padded = pad_sequences(test_seq, maxlen=max_length)

1 print("train : ", len(seq[0]))
2 print("test : ", len(test_seq[0]))

```

```

train : 118
test : 173

```

```

1 import nltk
2 nltk.download('punkt')
3 nltk.download('stopwords')

```

```

[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
True

```

```

1 import string
2 from nltk.tokenize import word_tokenize
3 from nltk.corpus import stopwords
4
5 output = []
6

```

```

7 for line in train_input:
8     tokens = word_tokenize(line)
9     tokens = [word.lower() for word in tokens]
10    processed = str.maketrans('', '', string.punctuation)
11    translated = [word.translate(processed) for word in tokens]
12    words = [word for word in translated if word.isalpha()]
13    stopWords = set(stopwords.words('english'))
14    words = [word for word in words if not word in stopWords]
15    output.append(words)

1 print(output[0])
2 print(len(output))

['b', 'absolutely', 'terrible', 'movie', 'nt', 'lured', 'christopher', 'walken',
25000

1 import gensim
2
3 model_1 = gensim.models.Word2Vec(sentences = output, size= embedding_dim, window
4 words = list(model_1.wv.vocab)

1 model_1.wv.save_word2vec_format("w2v_embedding.txt", binary = False)

1 import os
2 path_to_w2v_file = "w2v_embedding.txt"
3
4 embeddings_index = {}
5 with open(path_to_w2v_file) as f:
6     for line in f:
7         word, coefs = line.split(maxsplit=1)
8         coefs = np.fromstring(coefs, "f", sep=" ")
9         embeddings_index[word] = coefs
10
11 print("Found %s word vectors." % len(embeddings_index))

Found 93966 word vectors.

1 num_tokens = len(word_index) + 1
2 hits = 0
3 misses = 0
4
5 embedding_matrix = np.zeros((num_tokens, embedding_dim))
6
7 for word, i in word_index.items():
8     embedding_vector = embeddings_index.get(word)
9     if embedding_vector is not None:
10         # Words not found in embedding index will be all-zeros.
11         # This includes the representation for "padding" and "OOV"

```

```

12         embedding_matrix[i] = embedding_vector
13         hits += 1
14     else:
15         misses += 1
16 print("Converted %d words (%d misses)" % (hits, misses))
    Converted 70670 words (15869 misses)

1 from keras.layers import Embedding
2 from keras.initializers import Constant
3
4 embedding_layer = Embedding(num_tokens, embedding_dim, embeddings_initializer= Co

1 ## LSTM Model
2 from keras.models import Sequential
3 from keras.layers import Dense, LSTM, GRU
4 from keras.layers.embeddings import Embedding
5
6 model_1 = Sequential()
7 model_1.add(embedding_layer)
8 model_1.add(LSTM(units=64, dropout=0.2))
9 model_1.add(Dense(32, activation='relu'))
10 model_1.add(Dense(1, activation='sigmoid'))
11
12 model_1.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
embedding_1 (Embedding)	(None, 120, 100)	8654000
lstm (LSTM)	(None, 64)	42240
dense (Dense)	(None, 32)	2080
dense_1 (Dense)	(None, 1)	33
=====		
Total params: 8,698,353		
Trainable params: 44,353		
Non-trainable params: 8,654,000		
=====		

```

1 model_1.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy']
2 history = model_1.fit(padding, train_labels, epochs=20, validation_data = (test_x

Epoch 1/20
782/782 [=====] - 37s 38ms/step - loss: 0.5588 - accuracy: 0.7500
Epoch 2/20

```

```

782/782 [=====] - 29s 37ms/step - loss: 0.4431 - accura
Epoch 3/20
782/782 [=====] - 29s 37ms/step - loss: 0.4215 - accura
Epoch 4/20
782/782 [=====] - 29s 37ms/step - loss: 0.4035 - accura
Epoch 5/20
782/782 [=====] - 28s 36ms/step - loss: 0.3999 - accura
Epoch 6/20
782/782 [=====] - 28s 36ms/step - loss: 0.3898 - accura
Epoch 7/20
782/782 [=====] - 28s 36ms/step - loss: 0.3851 - accura
Epoch 8/20
782/782 [=====] - 28s 36ms/step - loss: 0.3772 - accura
Epoch 9/20
782/782 [=====] - 28s 36ms/step - loss: 0.3726 - accura
Epoch 10/20
782/782 [=====] - 28s 36ms/step - loss: 0.3632 - accura
Epoch 11/20
782/782 [=====] - 31s 40ms/step - loss: 0.3573 - accura
Epoch 12/20
782/782 [=====] - 28s 36ms/step - loss: 0.3478 - accura
Epoch 13/20
782/782 [=====] - 28s 36ms/step - loss: 0.3412 - accura
Epoch 14/20
782/782 [=====] - 28s 36ms/step - loss: 0.3316 - accura
Epoch 15/20
782/782 [=====] - 28s 36ms/step - loss: 0.3266 - accura
Epoch 16/20
782/782 [=====] - 28s 36ms/step - loss: 0.3170 - accura
Epoch 17/20
782/782 [=====] - 28s 36ms/step - loss: 0.3071 - accura
Epoch 18/20
782/782 [=====] - 28s 36ms/step - loss: 0.3018 - accura
Epoch 19/20
782/782 [=====] - 29s 37ms/step - loss: 0.2941 - accura
Epoch 20/20
782/782 [=====] - 28s 36ms/step - loss: 0.2828 - accura

```

```

1 from sklearn import metrics
2 from sklearn.metrics import precision_score, recall_score
3
4 predicted_1 = model_1.predict(test_padded)
5
6 pred_1 = np.zeros(len(predicted_1))
7 for i, score in enumerate(predicted_1):
8     if score > 0.5:
9         pred_1[i] = 1
10    else:
11        pred_1[i] = 0
12
13 print(metrics.classification_report(test_labels, pred_1))
14 print("Precision = ", precision_score(test_labels, pred_1))
15 print("Recall = ", recall_score(test_labels, pred_1))

```

	precision	recall	f1-score	support
0	0.82	0.84	0.83	12500
1	0.83	0.82	0.83	12500
accuracy			0.83	25000
macro avg	0.83	0.83	0.83	25000
weighted avg	0.83	0.83	0.83	25000

Precision = 0.8347416314592135

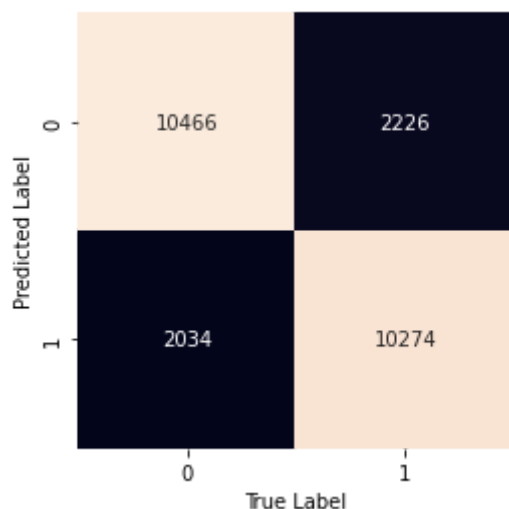
Recall = 0.82192

```

1 from sklearn.metrics import confusion_matrix
2 import seaborn as sns
3 import matplotlib.pyplot as plt
4
5 mat = confusion_matrix(test_labels, pred_1)
6 sns.heatmap(mat.T, square = True, annot = True, fmt = 'd', cbar = False)
7 plt.xlabel("True Label")
8 plt.ylabel("Predicted Label")

```

Text(91.68, 0.5, 'Predicted Label')



GLOVE

```

1 !wget http://nlp.stanford.edu/data/glove.6B.zip
2 !unzip -q glove.6B.zip

```

```

--2022-03-07 05:16:00-- http://nlp.stanford.edu/data/glove.6B.zip
Resolving nlp.stanford.edu (nlp.stanford.edu)... 171.64.67.140
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:80... connected
HTTP request sent, awaiting response... 302 Found
Location: https://nlp.stanford.edu/data/glove.6B.zip [following]
--2022-03-07 05:16:00-- https://nlp.stanford.edu/data/glove.6B.zip
Connecting to nlp.stanford.edu (nlp.stanford.edu)|171.64.67.140|:443... connecte
HTTP request sent, awaiting response... 301 Moved Permanently
Location: http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip [following]

```

```
--2022-03-07 05:16:00-- http://downloads.cs.stanford.edu/nlp/data/glove.6B.zip
Resolving downloads.cs.stanford.edu (downloads.cs.stanford.edu)... 171.64.64.22
Connecting to downloads.cs.stanford.edu (downloads.cs.stanford.edu)|171.64.64.22
HTTP request sent, awaiting response... 200 OK
Length: 862182613 (822M) [application/zip]
Saving to: 'glove.6B.zip'
```

```
glove.6B.zip          100%[=====>] 822.24M  5.11MB/s    in 2m 40s
```

```
2022-03-07 05:18:41 (5.13 MB/s) - 'glove.6B.zip' saved [862182613/862182613]
```

```
1 import os
2 path_to_glove_file = "glove.6B.100d.txt"
3
4 embeddings_index = {}
5 with open(path_to_glove_file) as f:
6     for line in f:
7         word, coefs = line.split(maxsplit=1)
8         coefs = np.fromstring(coefs, "f", sep=" ")
9         embeddings_index[word] = coefs
10
11 print("Found %s word vectors." % len(embeddings_index))
```

```
Found 400000 word vectors.
```

```
1 num_tokens = len(word_index) + 1
2 embedding_dim = 100 ## 100 dimensions
3 hits = 0 ## number of words that were found in the pretrained model
4 misses = 0 ## number of words that were missing in the pretrained model
5
6 # Prepare embedding matrix for our word list
7 embedding_matrix = np.zeros((num_tokens, embedding_dim))
8 for word, i in word_index.items():
9     embedding_vector = embeddings_index.get(word)
10    if embedding_vector is not None:
11        # Words not found in embedding index will be all-zeros.
12        # This includes the representation for "padding" and "OOV"
13        embedding_matrix[i] = embedding_vector
14        hits += 1
15    else:
16        misses += 1
17 print("Converted %d words (%d misses)" % (hits, misses))
```

```
Converted 60197 words (26342 misses)
```

```
1 from keras.layers import Embedding
2 from keras.initializers import Constant
3
4 embedding_layer = Embedding(num_tokens, embedding_dim, embeddings_initializer= Cc
```

```

1 ## LSTM Model
2 from keras.models import Sequential
3 from keras.layers import Dense, LSTM, GRU
4 from keras.layers.embeddings import Embedding
5
6 model_2 = Sequential()
7 model_2.add(embedding_layer)
8 model_2.add(LSTM(units=64, dropout=0.2))
9 model_2.add(Dense(32, activation='relu'))
10 model_2.add(Dense(1, activation='sigmoid'))
11
12 model_2.summary()

```

Model: "sequential_2"

Layer (type)	Output Shape	Param #
=====		
embedding_2 (Embedding)	(None, 120, 100)	8654000
lstm_2 (LSTM)	(None, 64)	42240
dense_4 (Dense)	(None, 32)	2080
dense_5 (Dense)	(None, 1)	33
=====		
Total params: 8,698,353		
Trainable params: 44,353		
Non-trainable params: 8,654,000		
=====		

```

1 model_2.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy']
2 history = model_2.fit(padding, train_labels, epochs=20, validation_data = (test_x

```

```

Epoch 1/20
782/782 [=====] - 31s 37ms/step - loss: 0.6761 - accuracy: 0.1250
Epoch 2/20
782/782 [=====] - 28s 36ms/step - loss: 0.4945 - accuracy: 0.2500
Epoch 3/20
782/782 [=====] - 28s 36ms/step - loss: 0.4385 - accuracy: 0.3750
Epoch 4/20
782/782 [=====] - 28s 36ms/step - loss: 0.4172 - accuracy: 0.5000
Epoch 5/20
782/782 [=====] - 28s 36ms/step - loss: 0.4015 - accuracy: 0.6250
Epoch 6/20
782/782 [=====] - 28s 36ms/step - loss: 0.3898 - accuracy: 0.7500
Epoch 7/20
782/782 [=====] - 28s 36ms/step - loss: 0.3760 - accuracy: 0.8750
Epoch 8/20
782/782 [=====] - 29s 37ms/step - loss: 0.3621 - accuracy: 0.9000
Epoch 9/20
782/782 [=====] - 28s 36ms/step - loss: 0.3513 - accuracy: 0.9250
Epoch 10/20

```



```

782/782 [=====] - 28s 36ms/step - loss: 0.3393 - accuracy: 0.8279
Epoch 11/20
782/782 [=====] - 28s 36ms/step - loss: 0.3271 - accuracy: 0.8381
Epoch 12/20
782/782 [=====] - 28s 36ms/step - loss: 0.3163 - accuracy: 0.8481
Epoch 13/20
782/782 [=====] - 29s 37ms/step - loss: 0.3029 - accuracy: 0.8581
Epoch 14/20
782/782 [=====] - 28s 36ms/step - loss: 0.2969 - accuracy: 0.8681
Epoch 15/20
782/782 [=====] - 29s 36ms/step - loss: 0.2831 - accuracy: 0.8781
Epoch 16/20
782/782 [=====] - 29s 37ms/step - loss: 0.2760 - accuracy: 0.8881
Epoch 17/20
782/782 [=====] - 28s 36ms/step - loss: 0.2648 - accuracy: 0.8981
Epoch 18/20
782/782 [=====] - 28s 36ms/step - loss: 0.2575 - accuracy: 0.9081
Epoch 19/20
782/782 [=====] - 28s 36ms/step - loss: 0.2511 - accuracy: 0.9181
Epoch 20/20
782/782 [=====] - 28s 36ms/step - loss: 0.2413 - accuracy: 0.9281

```

```

1 from sklearn import metrics
2 from sklearn.metrics import precision_score, recall_score
3
4 predicted_2 = model_2.predict(test_padded)
5
6 pred_2 = np.zeros(len(predicted_2))
7 for i, score in enumerate(predicted_2):
8     if score > 0.5:
9         pred_2[i] = 1
10    else:
11        pred_2[i] = 0
12
13 print(metrics.classification_report(test_labels, pred_2))
14 print("Precision = ", precision_score(test_labels, pred_2))
15 print("Recall = ", recall_score(test_labels, pred_2))

```

	precision	recall	f1-score	support
0	0.86	0.82	0.84	12500
1	0.83	0.87	0.85	12500
accuracy			0.84	25000
macro avg	0.85	0.84	0.84	25000
weighted avg	0.85	0.84	0.84	25000

```

Precision = 0.8279381521821921
Recall = 0.8696


```

```

1 from sklearn.metrics import confusion_matrix
2 import seaborn as sns

```

```
3 import matplotlib.pyplot as plt
4
5 mat = confusion_matrix(test_labels, pred_2)
6 sns.heatmap(mat.T, square = True, annot = True, fmt = 'd', cbar = False)
7 plt.xlabel("True Label")
8 plt.ylabel("Predicted Label")
9 plt.text(91.68, 0.5, 'Predicted Label')
```



! 0s completed at 7:50 AM

● ×